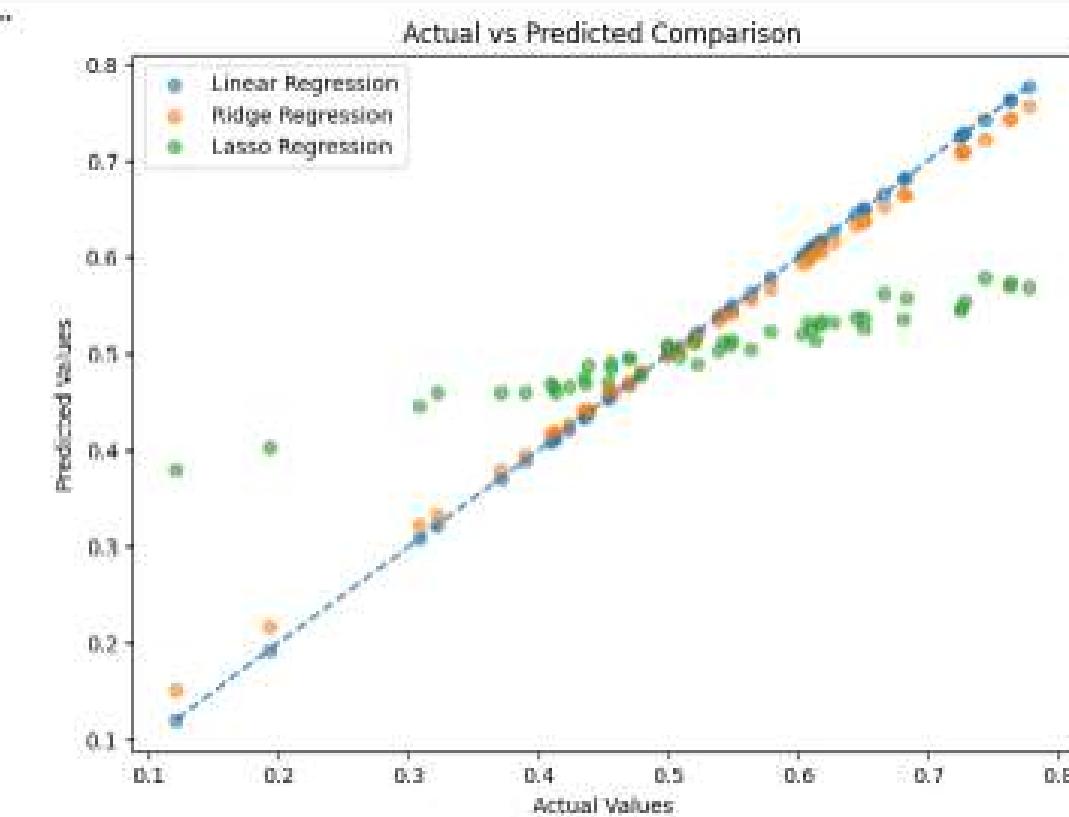


```
plt.figure(figsize=(8,6))

plt.scatter(y_test, y_pred_lr, label="Linear Regression", alpha=0.6)
plt.scatter(y_test, y_pred_ridge, label="Ridge Regression", alpha=0.6)
plt.scatter(y_test, y_pred_lasso, label="Lasso Regression", alpha=0.6)

plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()],
         linestyle='--')

plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs Predicted Comparison")
plt.legend()
plt.show()
```

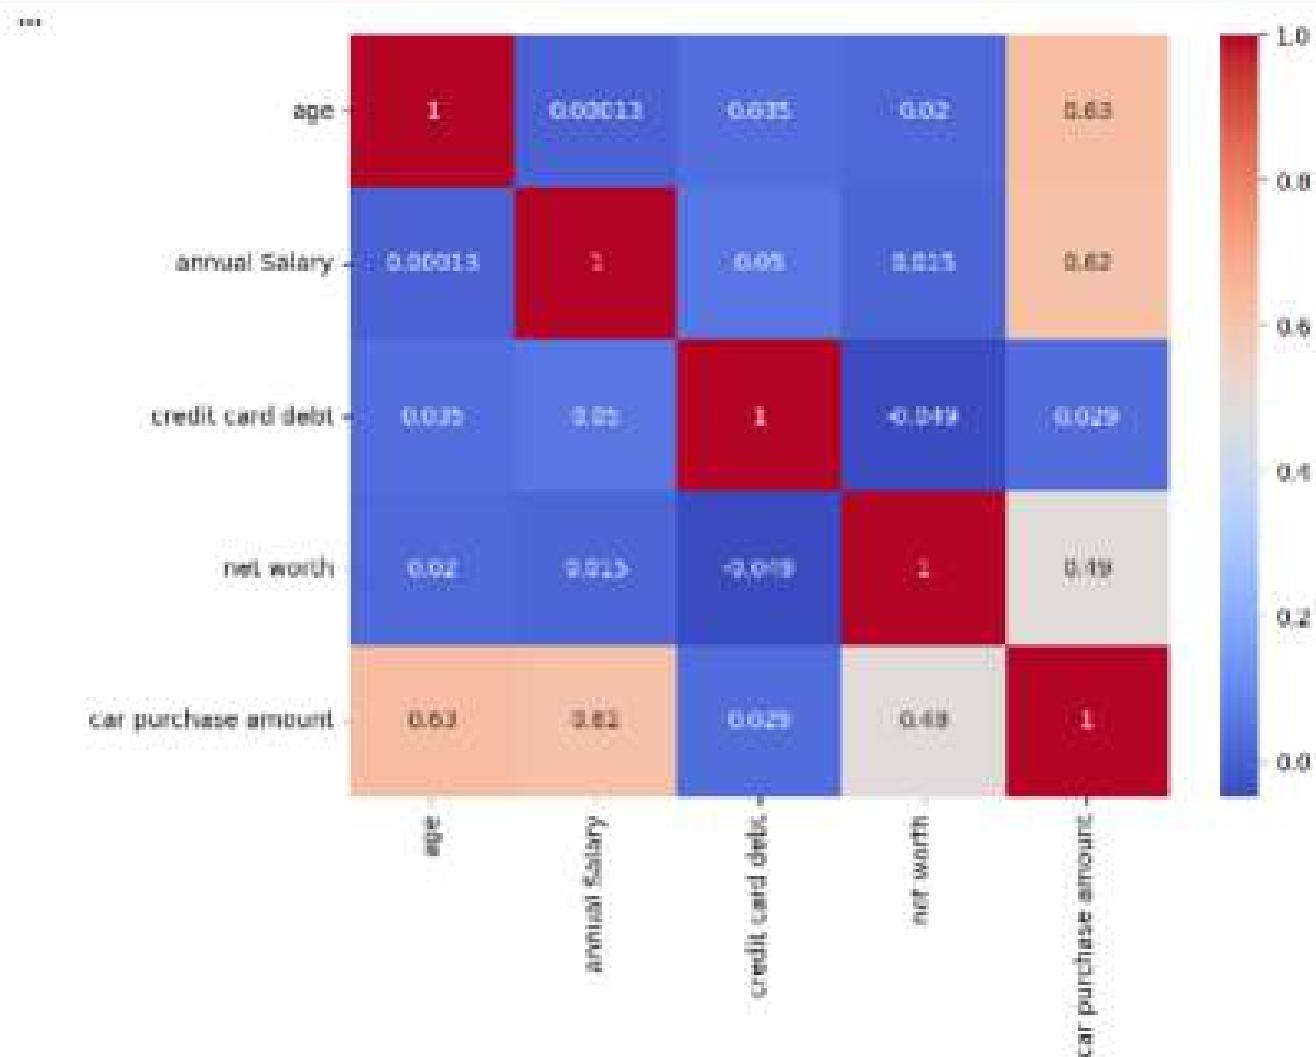


The plot shows that Linear and Ridge Regression predictions closely match the actual values, indicating high accuracy, while Lasso Regression predictions deviate more, reflecting lower performance due to feature shrinkage.

Age and annual salary show strong positive correlation with car purchase amount

Credit card debt has very low correlation, so it has minimal impact

```
❶ import seaborn as sns  
import matplotlib.pyplot as plt  
plt.figure(figsize=(8,5))  
sns.heatmap(df2.corr(), annot=True, cmap="coolwarm")  
plt.show()
```



```

# Import required libraries
import pandas as pd
from textblob import TextBlob

# Step 1: Ensure review_content is string (IMPORTANT)
df['review_content'] = df['review_content'].astype(str)

# Step 2: Apply sentiment analysis using TextBlob
df['sentiment'] = df['review_content'].apply(
    lambda text: TextBlob(text).sentiment.polarity
)

# Step 3: Sort reviews by sentiment score
# Most positive reviews
positive_reviews = df.sort_values(by='sentiment', ascending=False)

# Most negative reviews
negative_reviews = df.sort_values(by='sentiment', ascending=True)

# Step 4: Display top 10 positive and negative reviews
top_positive = positive_reviews[['product_id', 'user_id', 'review_content', 'sentiment']].head(10)
top_negative = negative_reviews[['product_id', 'user_id', 'review_content', 'sentiment']].head(10)

print("\nTop 10 Positive Reviews:")
print(top_positive)

print("\nTop 10 Negative Reviews:")
print(top_negative)

```

... Top 10 Positive Reviews:

	product_id	user_id	review_content	sentiment
1464	134	433	475	0.0
8	345	623	604	0.0
1	848	88	413	0.0
2	810	849	674	0.0
3	643	254	189	0.0
4	588	17	129	0.0
5	771	210	518	0.0
6	761	662	123	0.0
7	614	1162	1122	0.0
1448	952	24	1006	0.0

Top 10 Negative Reviews:

	product_id	user_id	review_content	sentiment
1455	53	82	5	-0.8
1454	785	134	782	-0.8
1453	348	1147	1052	-0.8
1452	1245	755	644	-0.8
1451	158	358	98	-0.8
1450	1314	717	888	-0.8
1449	1273	344	648	-0.8
1448	952	24	1006	-0.8
1447	298	63	144	-0.8
1446	984	893	554	-0.8

Answer 8: Discounted price and rating have a weak positive correlation. This means that products with higher discounted price have slightly higher ratings, but the relationship is not very strong.

```
# Linear Regression
lr = LinearRegression()
lr.fit(x_train, y_train)
y_pred_lr = lr.predict(x_test)

# Ridge Regression
ridge = Ridge(alpha=1.0)
ridge.fit(x_train, y_train)
y_pred_ridge = ridge.predict(x_test)

# Lasso Regression
lasso = Lasso(alpha=0.01)
lasso.fit(x_train, y_train)
y_pred_lasso = lasso.predict(x_test)
```

3. Evaluation Metrics Function

```
def evaluate_model(name, y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)

    return [name, mae, mse, rmse, r2]
```

4. Evaluate All Models

```
results = []

results.append(evaluate_model("Linear Regression", y_test, y_pred_lr))
results.append(evaluate_model("Ridge Regression", y_test, y_pred_ridge))
results.append(evaluate_model("Lasso Regression", y_test, y_pred_lasso))

results_df = pd.DataFrame(
    results,
    columns=["Model", "MAE", "MSE", "RMSE", "R2 Score"]
)

results_df
```

	Model	MAE	MSE	RMSE	R2 Score	
0	Linear Regression	0.000016	3.992213e-10	0.000020	1.000000	
1	Ridge Regression	0.009547	1.403898e-04	0.011849	0.999089	
2	Lasso Regression	0.085623	1.135407e-02	0.106555	0.441032	

```
[0.54146119],  
[0.38927052],  
[0.82661214],  
[0.40998855],  
[0.37072477],  
[0.46885649],  
[0.44642818],  
[0.53868386]))
```

Train the model

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()  
model.fit(x_train, y_train)  
...  
LinearRegression  
LinearRegression()
```

Make predictions

```
y_pred = model.predict(x_test)
```

Evaluate the model

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
print("MSE:", mean_squared_error(y_test, y_pred))  
print("R2 Score:", r2_score(y_test, y_pred))
```

```
MSE: 3.992213e-13  
R2 Score: 0.999999983468698
```

MSE (Mean Squared Error) It measures how far the predicted values (`y_pred`) are from the actual values (`y_test`).

Your value: 3.99e-10 → extremely small, almost 0.

Interpretation: The predictions are almost perfect, almost exactly equal to the actual values.

R2 Score

Measures how well the model explains the variance in the target.

Range: 0 to 1 (sometimes negative if very bad).

Your value: 0.99999998 → almost 1, meaning the model explains nearly 100% of the variance.